



Chapter 5: Advanced SQL

Instructor : Fatemeh Mansoori

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Functions and Procedures



Functions and Procedures

- Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements.
- These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, or C++.
- The syntax we present here is defined by the SQL standard.
 - Most databases implement nonstandard versions of this syntax.



Define function in Postgresql

```
CREATE [OR REPLACE] FUNCTION functionName (someParameter 'parameterType')
RETURNS 'DATATYPE'
AS $_block_name_$
DECLARE
    --declare something
BEGIN
    --do something
    --return something
END;
$_block_name_$
LANGUAGE plpgsql;
```



Functions

Create a function that get the department name and returns the number of instructors in that department

```
create or replace function instructor_count(dept_name varchar(20))
returns integer
as
$$
declare
icount integer;
begin
select count(*) into icount
from instructor as i
where i.dept_name = instructor_count.dept_name;
return icount;
end;
$$
language plpgsql;
```



Table Functions

- The SQL standard supports functions that can return tables as results; such functions are called **table functions**
- Example: Return all instructors in a given department

```
create or replace function instructor_of(dn varchar(20))
returns table(id varchar(5), name varchar(20), dept_name varchar(20), salary numeric(8,2)) as
$$
begin
return query
(select * from instructor as i
 where i.dept_name = dn
);
end;
$$
language plpgsql;
```

- Usage

```
select *
from table (instructor_of('Music'))
```



SQL Procedures

- The *dept_count* function could instead be written as procedure:
- In Postgresql :

```
create or replace procedure dept_count(in dn varchar(20), out icount integer)
as
$$
begin
    select count(*) into icount
    from instructor
    where dept_name=dn;
end;
$$
language plpgsql;
```

- The keywords **in** and **out** are parameters that are expected to have values assigned to them and parameters whose values are set in the procedure in order to return results.
- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.
call dept_count_proc('Physics', 0);



Example procedure

- Registers student after ensuring classroom capacity is not exceeded

```
create or replace procedure registration(id student.id%type,  
                                         course_id course.course_id%type,  
                                         sec_id section.sec_id%type,  
                                         semester section.semester%type,  
                                         year section.year%type) as  
  
$$  
declare r_count int; lmt int;  
begin  
    select count(*) into r_count  
    from takes t  
    where t.course_id = registration.course_id and t.sec_id = registration.sec_id and  
    t.semester = registration.semester and t.year = registration.year;  
    select capacity into lmt  
    from section s natural join classroom c  
    where s.course_id = registration.course_id and s.sec_id = registration.sec_id and  
    s.semester = registration.semester and s.year = registration.year;  
    if r_count < lmt then  
        insert into takes(id, course_id, sec_id, semester, year)  
        values (registration.id, registration.course_id, registration.sec_id,  
                registration.semester, registration.year);  
    else  
        raise 'Capacity is full';  
    end if;  
end;  
$$  
language plpgsql;
```




Exception conditions

- Signaling of exception conditions, and declaring handlers for exceptions

```
declare out_of_classroom_seats condition
declare exit handler for out_of_classroom_seats
begin
...
end
```
- The statements between the **begin** and the **end** can raise an exception by executing “**signal** *out_of_classroom_seats*”
- The handler says that if the condition arises the action to be taken is to exit the enclosing the **begin end** statement.



SQL Procedures (Cont.)

- Procedures and functions can be invoked also from dynamic SQL
- SQL allows more than one procedure of the so long as the number of arguments of the procedures with the same name is different.
- The name, along with the number of arguments, is used to identify the procedure.



Language Constructs for Procedures & Functions

- SQL supports constructs that gives it almost all the power of a general-purpose programming language.
 - Warning: most database systems implement their own variant of the standard syntax below.
- **Postgresql language constructs :**
- Conditional statements (**if-then-else**)

```
if boolean expression then  
    statement (s)  
elseif boolean expression then  
    statement(s)  
else  
    statement(s)  
end if
```



Language Constructs (Cont.)

```
IF number = 0 THEN
    result := 'zero';
ELSIF number > 0 THEN
    result := 'positive';
ELSIF number < 0 THEN
    result := 'negative';
ELSE
    -- hmm, the only other possibility is that number is null
    result := 'NULL';
END IF;
```



Language Constructs (Cont.)

```
CASE x
  WHEN 1, 2 THEN
    msg := 'one or two';
  ELSE
    msg := 'other value than one or two';
END CASE;
```

```
CASE
  WHEN x BETWEEN 0 AND 10 THEN
    msg := 'value is between zero and ten';
  WHEN x BETWEEN 11 AND 20 THEN
    msg := 'value is between eleven and twenty';
END CASE;
```



Language Constructs (Cont.)

- Loops:

```
LOOP
  -- some computations
  IF count > 0 THEN
    EXIT;  -- exit loop
  END IF;
END LOOP;
```

```
LOOP
  -- some computations
  EXIT WHEN count > 0;  -- same result as previous example
END LOOP;
```



Language Constructs (Cont.)

```
<<ablock>>  
BEGIN  
    -- some computations  
    IF stocks > 100000 THEN  
        EXIT ablock; -- causes exit from the BEGIN block  
    END IF;  
    -- computations here will be skipped when stocks > 100000  
END;
```

```
LOOP  
    -- some computations  
    EXIT WHEN count > 100;  
    CONTINUE WHEN count < 50;  
    -- some computations for count IN [50 .. 100]  
END LOOP;
```



Language Constructs (Cont.)

```
WHILE amount_owed > 0 AND gift_certificate_balance > 0 LOOP
    -- some computations here
END LOOP;
```

```
WHILE NOT done LOOP
    -- some computations here
END LOOP;
```

```
FOR i IN 1..10 LOOP
    -- i will take on the values 1,2,3,4,5,6,7,8,9,10 within the loop
END LOOP;
```

```
FOR i IN REVERSE 10..1 LOOP
    -- i will take on the values 10,9,8,7,6,5,4,3,2,1 within the loop
END LOOP;
```

```
FOR i IN REVERSE 10..1 BY 2 LOOP
    -- i will take on the values 10,8,6,4,2 within the loop
END LOOP;
```




Language Constructs (Cont.)

```
DO $$
DECLARE
city_names varchar;
BEGIN
FOR city_names IN SELECT city_name FROM major_cities
LOOP
RAISE NOTICE '%', city_names;
END LOOP;
END$$;
```

```
DO $$
DECLARE
emp_name record;
BEGIN
FOR emp_name IN SELECT first_name, last_name FROM employee LIMIT 10
LOOP
RAISE NOTICE '% %', emp_name.first_name, emp_name.last_name;
END LOOP;
END$$;
```



Language Constructs (Cont.)

```
FOR r IN SELECT * FROM foo
  WHERE fooid > 0
  LOOP
    -- can do some processing here
    RETURN NEXT r; -- return current row of SELECT
  END LOOP;
```



Triggers



Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
 - Specify the conditions under which the trigger is to be executed.
 - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
 - Syntax illustrated here may not work exactly on your database system; check the system manuals



Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
 - For example, **after update of *takes* on *grade***
- Values of attributes before and after an update can be referenced
 - **referencing old row as** : for deletes and updates
 - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
    when (nrow.grade = ' ')  
    begin atomic  
        set nrow.grade = null;  
    end;
```



Trigger in postgresql

```
CREATE TRIGGER trigger_name
    {BEFORE | AFTER} { event }
    ON table_name
    [FOR [EACH] { ROW | STATEMENT }]
    EXECUTE PROCEDURE trigger_function
```

- The event can be Insert, delete, update or truncate
- A row-level trigger is fired for each row while a statement-level trigger is fired for each transaction
- Suppose a table has 100 rows and two triggers that will be fired when a DELETE event occurs.
- If the DELETE statement deletes 100 rows, the row-level trigger will fire 100 times, once for each deleted row. On the other hand, a statement-level trigger will be fired for one time regardless of how many rows are deleted.



trigger function syntax

```
CREATE FUNCTION trigger_function()  
    RETURNS TRIGGER  
    LANGUAGE PLPGSQL  
AS $$  
BEGIN  
    -- trigger logic  
END;  
$$
```



Statement Level Triggers

- Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - Use **for each statement** instead of **for each row**
 - Use **referencing old table** or **referencing new table** to refer to temporary tables (called *transition tables*) containing the affected rows
 - Can be more efficient when dealing with SQL statements that update a large number of rows



When Not To Use Triggers

- Triggers were used earlier for tasks such as
 - Maintaining summary data (e.g., total salary of each department)
 - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
 - Databases today provide built in materialized view facilities to maintain summary data
 - Databases provide built-in support for replication



When Not To Use Triggers (Cont.)

- Risk of unintended execution of triggers, for example, when
 - Loading data from a backup copy
 - Replicating updates at a remote site
 - Trigger execution can be disabled before such actions.
- Other risks with triggers:
 - Error leading to failure of critical transactions that set off the trigger
 - Cascading execution



- Disable a trigger:
- when you disable a trigger, the trigger still exists in the database. However, the disabled trigger will not fire when an event associated with the trigger occurs

```
ALTER TABLE table_name  
DISABLE TRIGGER trigger_name | ALL
```

- Enable a trigger:

```
ALTER TABLE table_name  
ENABLE TRIGGER trigger_name | ALL;
```

- Delete a trigger :

```
DROP TRIGGER [IF EXISTS] trigger_name  
ON table_name [ CASCADE | RESTRICT ];
```



Recursive Queries



Recursion in SQL

- SQL:1999 permits recursive view definition
- Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
with recursive rec_prereq(course_id, prereq_id) as (  
    select course_id, prereq_id  
    from prereq  
    union  
    select rec_prereq.course_id, prereq.prereq_id,  
    from rec_rereq, prereq  
    where rec_prereq.prereq_id = prereq.course_id  
    )  
select *  
from rec_prereq;
```

This example view, *rec_prereq*, is called the *transitive closure* of the *prereq* relation



The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
 - This can give only a fixed number of levels of managers
 - Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work



The Power of Recursion

- Computing transitive closure using iteration, adding successive tuples to *rec_prereq*
 - The next slide shows a *prereq* relation
 - Each step of the iterative process constructs an extended version of *rec_prereq* from its recursive definition.
 - The final result is called the *fixed point* of the recursive view definition.



Advanced Aggregation Features



Ranking

- Ranking is done in conjunction with an order by specification.
- Suppose we are given a relation
student_grades(*ID*, *GPA*)
giving the grade-point average of each student
- Find the rank of each student.
- **select *ID*, rank() over (order by *GPA* desc) as *s_rank***
from *student_grades*
- An extra **order by** clause is needed to get them in sorted order
select *ID*, rank() over (order by *GPA* desc) as *s_rank*
from *student_grades*
order by *s_rank*
- Ranking may leave gaps: e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
 - **dense_rank** does not leave gaps, so next dense rank would be 2



Ranking

- Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)  
                    from student_grades B  
                    where B.GPA > A.GPA)) as s_rank  
from student_grades A  
order by s_rank,
```



Ranking (Cont.)

- Ranking can be done within partition of the data.
- “Find the rank of students within each department.”

```
select ID, dept_name,  
       rank () over (partition by dept_name order by GPA desc)  
       as dept_rank  
from dept_grades  
order by dept_name, dept_rank;
```
- Multiple **rank** clauses can occur in a single **select** clause.
- Ranking is done *after* applying **group by** clause/aggregation
- Can be used to find top-n results
 - More general than the **limit** *n* clause supported by many databases, since it allows top-n within each partition



Ranking (Cont.)

- Other ranking functions:
 - **percent_rank** (within partition, if partitioning is done)
 - **cume_dist** (cumulative distribution)
 - fraction of tuples with preceding values
 - **row_number** (non-deterministic in presence of duplicates)
- SQL:1999 permits the user to specify **nulls first** or **nulls last**
select *ID*,
 rank () over (order by *GPA* **desc nulls last)** **as** *s_rank*
from *student_grades*



Ranking (Cont.)

- For a given constant n , the ranking the function $ntile(n)$ takes the tuples in each partition in the specified order, and divides them into n buckets with equal numbers of tuples.
- E.g.,
select ID, ntile(4) over (order by GPA desc) as quartile
from student_grades;



Windowing

- Used to smooth out random variations.
- E.g., **moving average**: “Given sales values for each date, calculate for each date the average of the sales on that day, the previous day, and the next day”
- **Window specification** in SQL:
 - Given relation *sales(date, value)*
select *date*, **sum**(*value*) **over**
 (**order by** *date* **between rows** 1 **preceding** and 1 **following**)
from *sales*

Postgres syntax:



Windowing

- Examples of other window specifications:
 - **between rows unbounded preceding and current**
 - **rows unbounded preceding**
 - **range between 10 preceding and current row**
 - All rows with values between current row value –10 to current value
 - **range interval 10 day preceding**
 - Not including current row



Windowing (Cont.)

- Can do windowing within partitions
- E.g., Given a relation *transaction* (*account_number*, *date_time*, *value*), where value is positive for a deposit and negative for a withdrawal
 - “Find total balance of each account after each transaction on the account”

```
select account_number, date_time,  
       sum (value) over  
         (partition by account_number  
          order by date_time  
          rows unbounded preceding)  
       as balance  
from transaction  
order by account_number, date_time
```




Windowing (Cont.)

- Postgresql syntax

```
SELECT
    product_name,
    group_name,
    price,
    LAST_VALUE (price) OVER (
        PARTITION BY group_name
        ORDER BY
            price RANGE BETWEEN UNBOUNDED PRECEDING
            AND UNBOUNDED FOLLOWING
    ) AS highest_price_per_group
FROM
    products
```



End of Chapter 5